

TEAM 13

FINAL PROJECT

CPE 313

Advanced Machine Learning
and Deep Learning

Dejoras, Dylan James Villamor, Kurt Russel

**NLP AND STT SYSTEMS STRUGGLE
WITH SENTIMENT DETECTION DUE
TO BIASES, NOISE, AND
CHALLENGES IN PROCESSING
SPONTANEOUS SPEECH AND LOW-
RESOURCE LANGUAGES.**



SOURCE:

[HTTPS://WWW.RESEARCHGATE.NET/POST/WHAT_ARE_THE_CHALLENGES_RESEARCHERS_MAY_FACE_DURING_LINGUISTIC_TEXT_ANALYSIS_HOW_CAN_THESE_CHALLENGES_BE_OVERCOME](https://www.researchgate.net/post/What_are_the_challenges_researchers_may_face_during_linguistic_text_analysis_how_can_these_challenges_be_overcome)

Speech-to-Text and Tonal Analysis for

VOICE-BASED

SENTIMENT

ANALYSIS

using Small Language Modeling

Dejoras, Dylan James Villamor, Kurt Russel

INTRODUCTION

Sentiment analysis helps industries improve user experience by using both text and voice data, with voice-based analysis capturing emotions through tone, pitch, and speech speed.

OBJECTIVES

- **Compare sentiment classification using text-based and audio-based inputs.**
- **Evaluate practical performance of each input type in classification tasks.**
- **Explore deep learning models for audio-based and text-based sentiment recognition.**

INTRODUCING LANGUAGE MODELS

Bidirectional Encoder
Representations from
Transformers

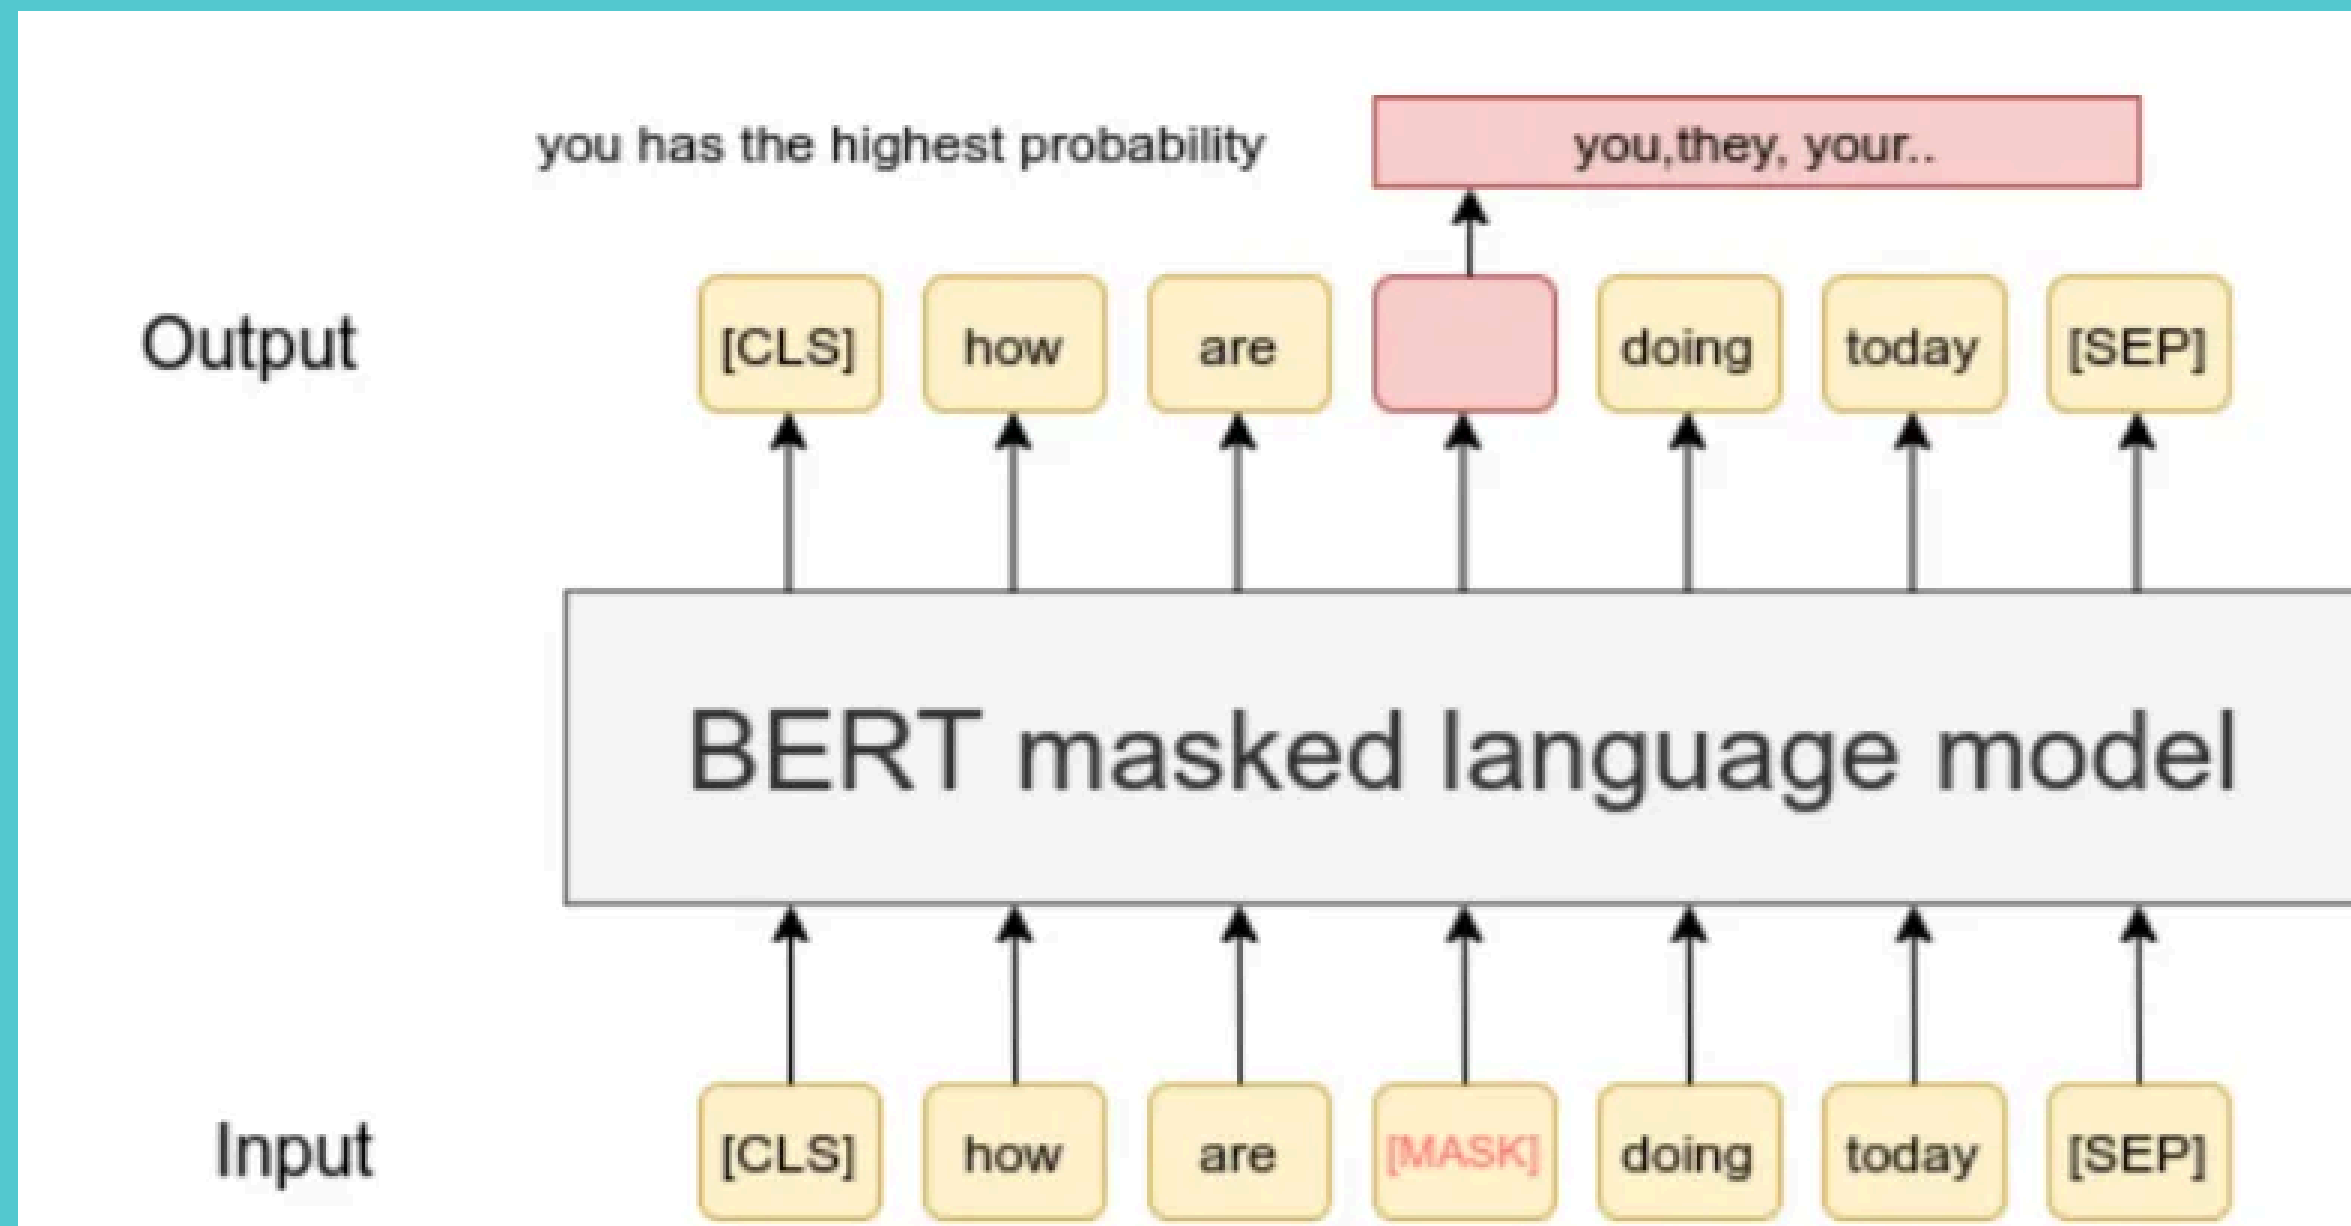
&

Efficiently Learning an
Encoder that
Classifies Token
Replacements
Accurately

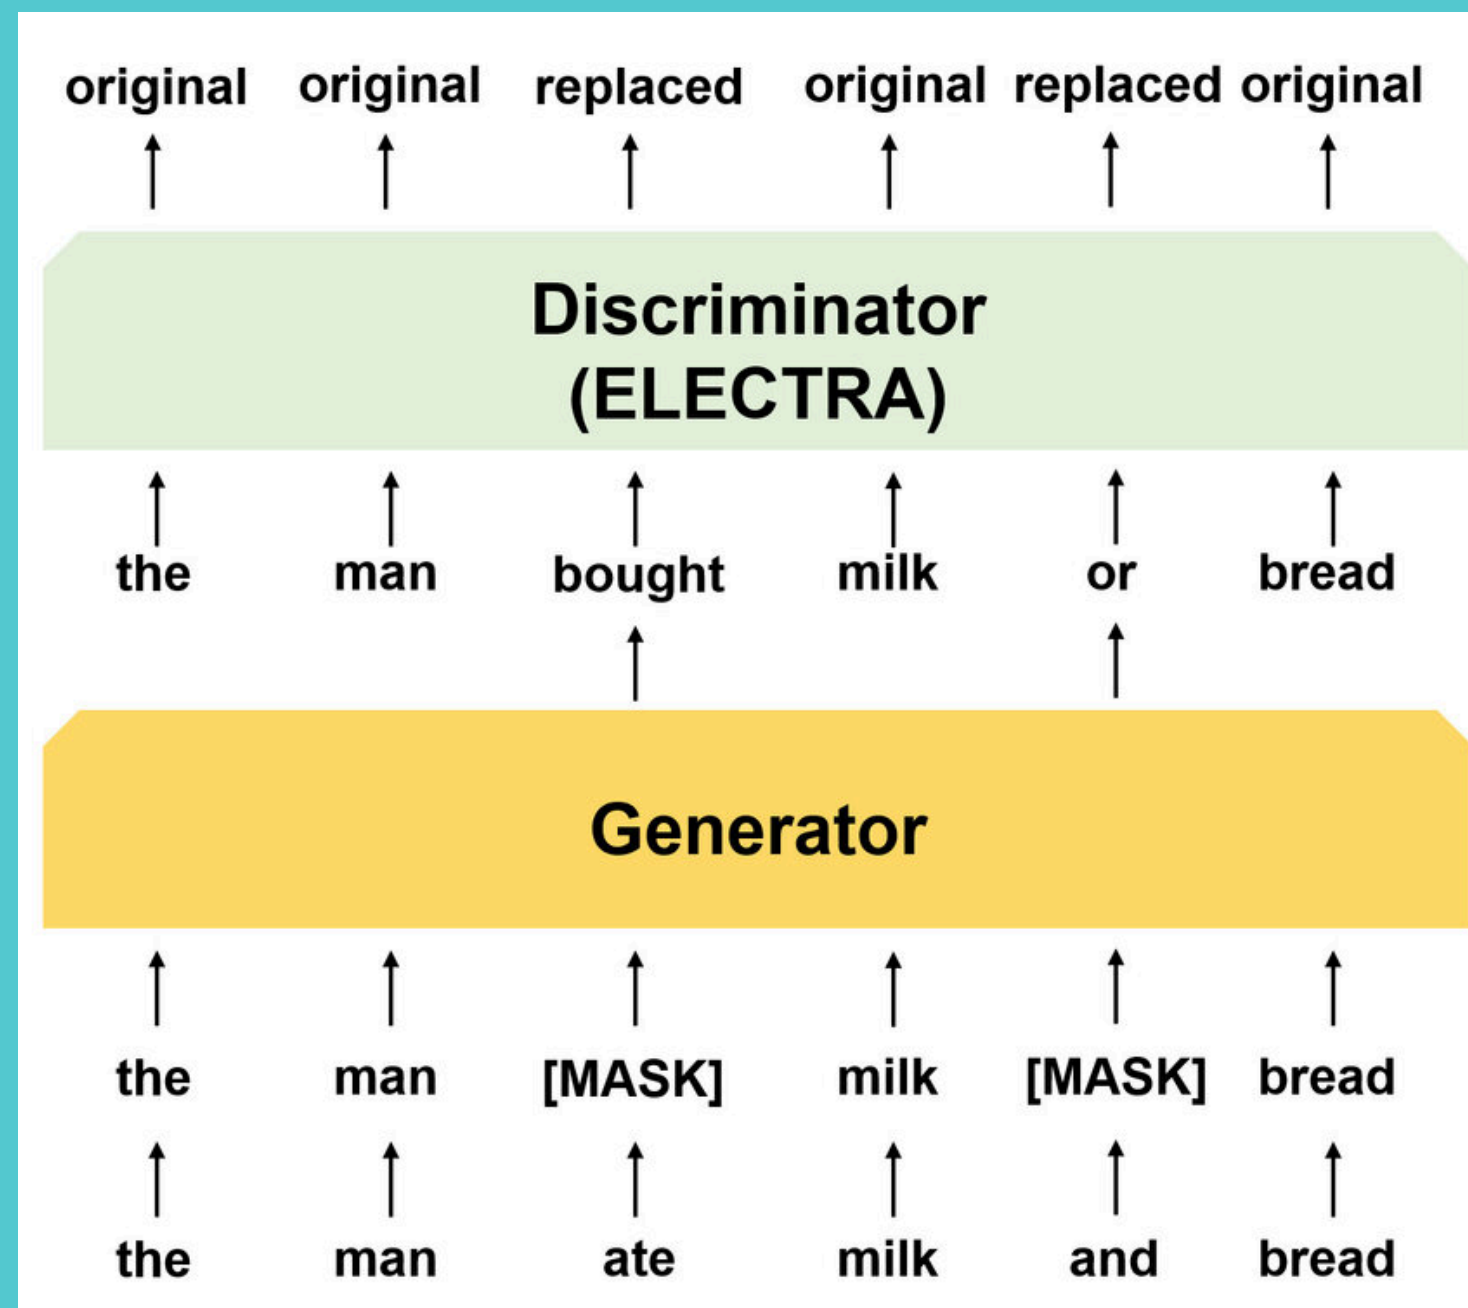
BERT

ELECTRA

BERT



ELECTRA



MODEL VARIANTS

RoBERTa

ELECTRA-
Small

TinyBERT

DATASETS



LAKSHMIPATHI N · UPDATED 6 YEARS AGO

IMDB Dataset of 50K Movie Reviews

Large Movie Review Dataset



ABHISHEK SHRIVASTAVA · UPDATED 4 YEARS AGO

Sentiment Analysis Dataset

Sentiment Analysis Dataset



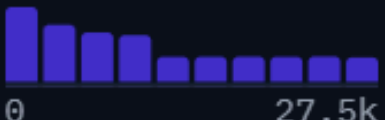



PASSIONATE-NLP · UPDATED 4 YEARS AGO

Twitter Sentiment Analysis

Entity-level sentiment analysis on multi-lingual tweets.

DATASETS

Datasets: Sp1786/multiclass-sentiment-analysis-dataset

| id | text | label | sentiment |
|---|--|---|---|
| int64 | string · lengths | int64 | string · classes |
|  |  |  |  |
| 0 27.5k | 1 2.18k | 0 2 | 3 values |
| 9,536 | Cooking microwave pizzas, yummy | 2 | positive |
| 6,135 | Any plans of allowing sub tasks to show up in the widget? | 1 | neutral |
| 17,697 | I love the humor, I just reworded it. Like saying 'group therapy' instead'a 'gang banging'. Keeps... | 2 | positive |
| 14,182 | naw idk what ur talkin about | 1 | neutral |
| 17,840 | That sucks to hear. I hate days like that | 0 | negative |
| 3,655 | Umm yeah. That's probably a pretty good note to self because eeeeeewwwwwwww. | 2 | positive |
| 719 | whatever do you mean? | 1 | neutral |

< Previous

1

2

3

...

313










Next >

118,334
samples

MODEL VARIANTS

| | MobileBERT | TinyBERT | ELECTRA-Small |
|---------------|--------------|--------------|---------------|
| ACCURACY | 89% | 81% | 89% |
| SIZE | 96.5 mb | 17.15 mb | 53.01 mb |
| TRAINING TIME | 9hrs. 20min. | 1hrs. 47min. | 2hrs. 4min. |

MODEL VARIANTS


| | MobileBERT | TinyBERT | ELECTRA-Small |
|---------------|---|---|---|
| ACCURACY |  |  |  |
| SIZE |  |  |  |
| TRAINING TIME |  |  |  |

THE CHALLENGE OF LIMITED RESOURCES

As large language models deliver high success rates, developers have turned to training models tailored to their own needs in this field. However, since training language models requires hardware capable of performing complex and numerous calculations quickly, it consumes a lot of energy.

SOURCE: Z. ÖRPEK, B. TURAL AND Z. DESTAN, "THE LANGUAGE MODEL REVOLUTION: LLM AND SLM ANALYSIS," 2024 8TH INTERNATIONAL ARTIFICIAL INTELLIGENCE AND DATA PROCESSING SYMPOSIUM (IDAP), MALATYA, TURKIYE, 2024, PP. 1-4, DOI: 10.1109/IDAP64064.2024.10710677.

LARGE LANGUAGE MODELS IN SOME CASES




SOURCE: [HTTPS://RESEARCH.AIMULTIPLE.COM/SENTIMENT-ANALYSIS-BENCHMARK/#BENCHMARK-DATASET-AND-METHODOLOGY](https://research.aimultiple.com/sentiment-analysis-benchmark/#benchmark-dataset-and-methodology)

PROCESS




1

SPEECH AS
INPUT



2

TEXT TRANSCRIPT USING
SPEECH-TO-TEXT MODEL



3

CLASSIFIED SENTIMENT
(POSITIVE, NEUTRAL, NEGATIVE)



VOICE-BASED

SENTIMENT ANALYSIS

DATASETS

RAVDESS (Ryerson Audio-Visual Database of Emotional Speech and Song)

[✎ Edit](#)

The Ryerson Audio-Visual Database of Emotional Speech and Song (RAVDESS) contains 7,356 files (total size: 24.8 GB). The database contains 24 professional actors (12 female, 12 male), vocalizing two lexically-matched statements in a neutral North American accent. Speech includes calm, happy, sad, angry, fearful, surprise, and disgust expressions, and song contains calm, happy, sad, angry, and fearful emotions. Each expression is produced at two levels of emotional intensity (normal, strong), with an additional neutral expression. All conditions are available in three modality formats: Audio-only (16bit, 48kHz .wav), Audio-Video (720p H.264, AAC 48kHz, .mp4), and Video-only (no sound). Note, there are no song files for Actor_18.

Usage 



RAVDESS

**'NEUTRAL', 'CALM', 'HAPPY', 'SAD', 'ANGRY',
'FEARFUL', 'DISGUST', 'SURPRISED'
2304 SAMPLES (RESAMPLED)**

TESS

**'NEUTRAL', 'CALM', 'HAPPY', 'SAD', 'ANGRY',
'FEARFUL', 'DISGUST', 'SURPRISED'
4800 SAMPLES (RESAMPLED)**

Toronto emotional speech set (TESS)

Permanent URI for this collection <https://hdl.handle.net/1807/24487>

These stimuli were modeled on the Northwestern University Auditory Test No. 6 (NU-6; Tillman & Carhart, 1966). A set of 200 target words were spoken in the carrier phrase "Say the word ____" by two actresses (aged 26 and 64 years) and recordings were made of the set portraying each of seven emotions (anger, disgust, fear, happiness, pleasant surprise, sadness, and neutral). There are 2800 stimuli in total. Two actresses were recruited from the Toronto area. Both actresses speak English as their first language, are university educated, and have musical training. Audiometric testing indicated that both actresses have thresholds within the normal range.


Authors: Kate Dupuis, M. Kathleen Pichora-Fuller

University of Toronto, Psychology Department, 2010.

Files are also available to download via [Scholars Portal Dataverse](https://doi.org/10.5683/SP2/E8H2MF). <https://doi.org/10.5683/SP2/E8H2MF>

This collection is published under Creative Commons license Attribution-NonCommercial-NoDerivatives 4.0 International.

DATASETS



Audio

SAVEE (Surrey Audio-Visual Expressed Emotion)

The **Surrey Audio-Visual Expressed Emotion (SAVEE)** dataset was recorded as a pre-requisite for the development of an automatic emotion recognition system. The database consists of recordings from 4 male actors in 7 different emotions, 480 British English utterances in total. The sentences were chosen from the standard TIMIT corpus and phonetically-balanced for each emotion. The data were recorded in a visual media lab with high quality audio-visual equipment, processed and labeled. To check the quality of performance, the recordings were evaluated by 10 subjects under audio, visual and audio-visual conditions. Classification systems were built using standard features and classifiers for each of the audio, visual and audio-visual modalities, and speaker-independent recognition rates of 61%, 65% and 84% achieved respectively.

SAVEE

**'NEUTRAL', 'CALM', 'HAPPY', 'SAD', 'ANGRY',
'FEARFUL', 'DISGUST', 'SURPRISED'**

720 SAMPLES (RESAMPLED)

CREMA-D

**'NEUTRAL', 'CALM', 'HAPPY', 'SAD', 'ANGRY',
'FEARFUL', 'DISGUST', 'SURPRISED'**

15252 SAMPLES (RESAMPLED)

CREMA-D

CREMA-D is an emotional multimodal actor data set of 7,442 original clips from 91 actors. These clips were from 48 male and 43 female actors between the ages of 20 and 74 coming from a variety of races and ethnicities (African America, Asian, Caucasian, Hispanic, and Unspecified).

Actors spoke from a selection of 12 sentences. The sentences were presented using one of six different emotions (Anger, Disgust, Fear, Happy, Neutral, and Sad) and four different emotion levels (Low, Medium, High, and Unspecified).

Participants rated the emotion and emotion levels based on the combined audiovisual presentation, the video alone, and the audio alone. Due to the large number of ratings needed, this effort was crowd-sourced and a total of 2443 participants each rated 90 unique clips, 30 audio, 30 visual, and 30 audio-visual. 95% of the clips have more than 7 ratings.

DATASETS



HAPPY, SURPRISED





SAD, ANGRY, FEARFUL, DISGUST, FEAR

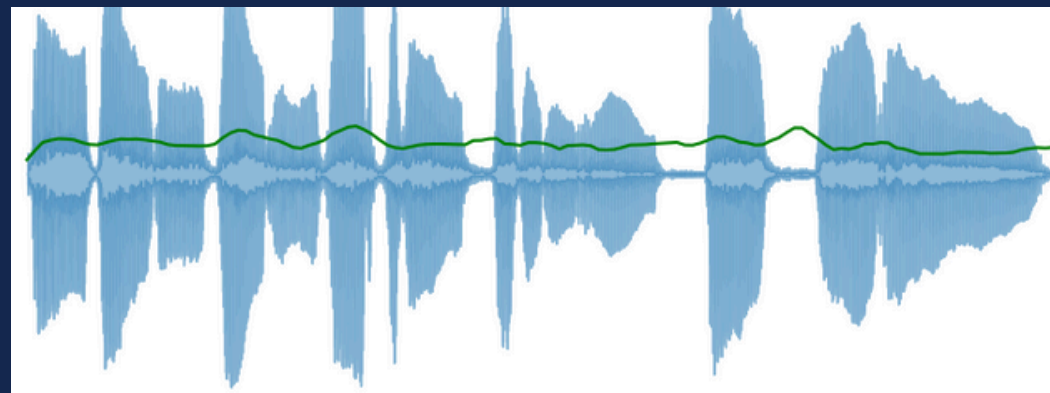


NEUTRAL


CNN (PREP)



ZERO CROSSING RATE



ROOT MEAN SQUARED



**MEL-FREQUENCY CEPSTRAL
COEFFICIENTS**

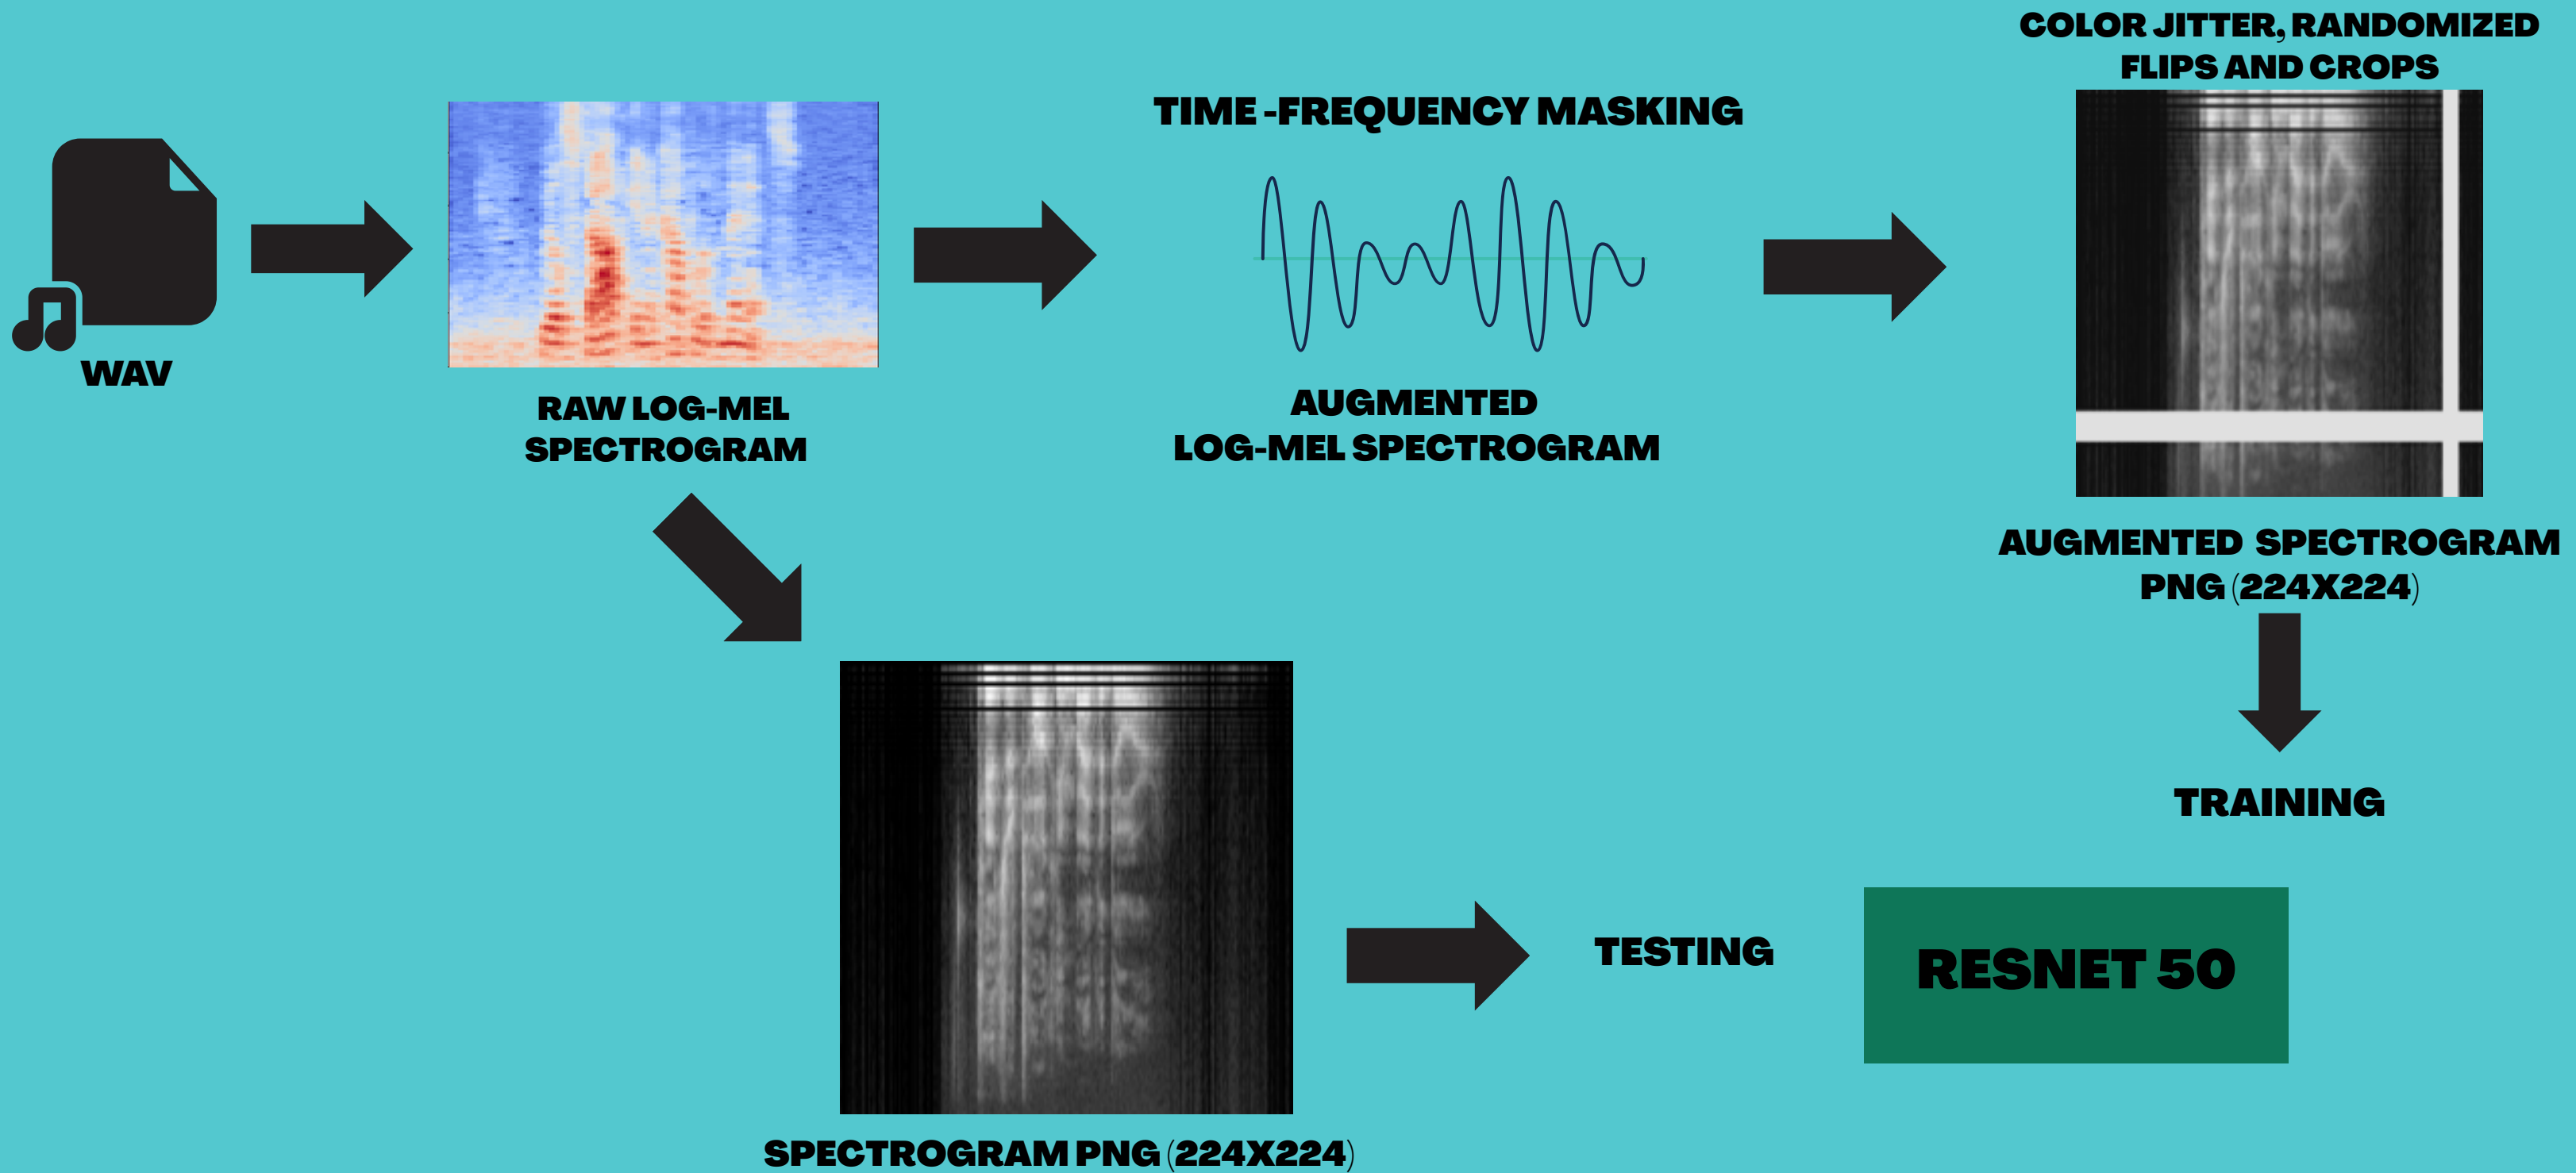
CNN

| Layer Block | Description | Output Shape (Example) |
|-------------------------------|---|------------------------|
| Conv1 + BN + Pool (x2) | Conv1d layers with BatchNorm and MaxPool reduce time dimension by 4× | [B, 512, T/4] |
| Dropout | Regularization (p=0.2) | [B, 512, T/4] |
| Conv1 + BN + Pool (x3) | Conv1d layers with BatchNorm and MaxPool further reduce time dimension by 8× | [B, 128, T/32] |
| Dropout | Regularization (p=0.2) | [B, 128, T/32] |


CNN

| Layer Block | Description | Output Shape (Example) |
|-------------------|---|------------------------|
| Flatten + FC + BN | Flatten features and fully connected layer to 512 units | [B, 512] |
| FC (Output) | Final fully connected layer to 3 output logits | [B, 3] |

RESNET 50 (PREP)



RESNET 50



WAV2VEC 2.0 (PREP)



WAV

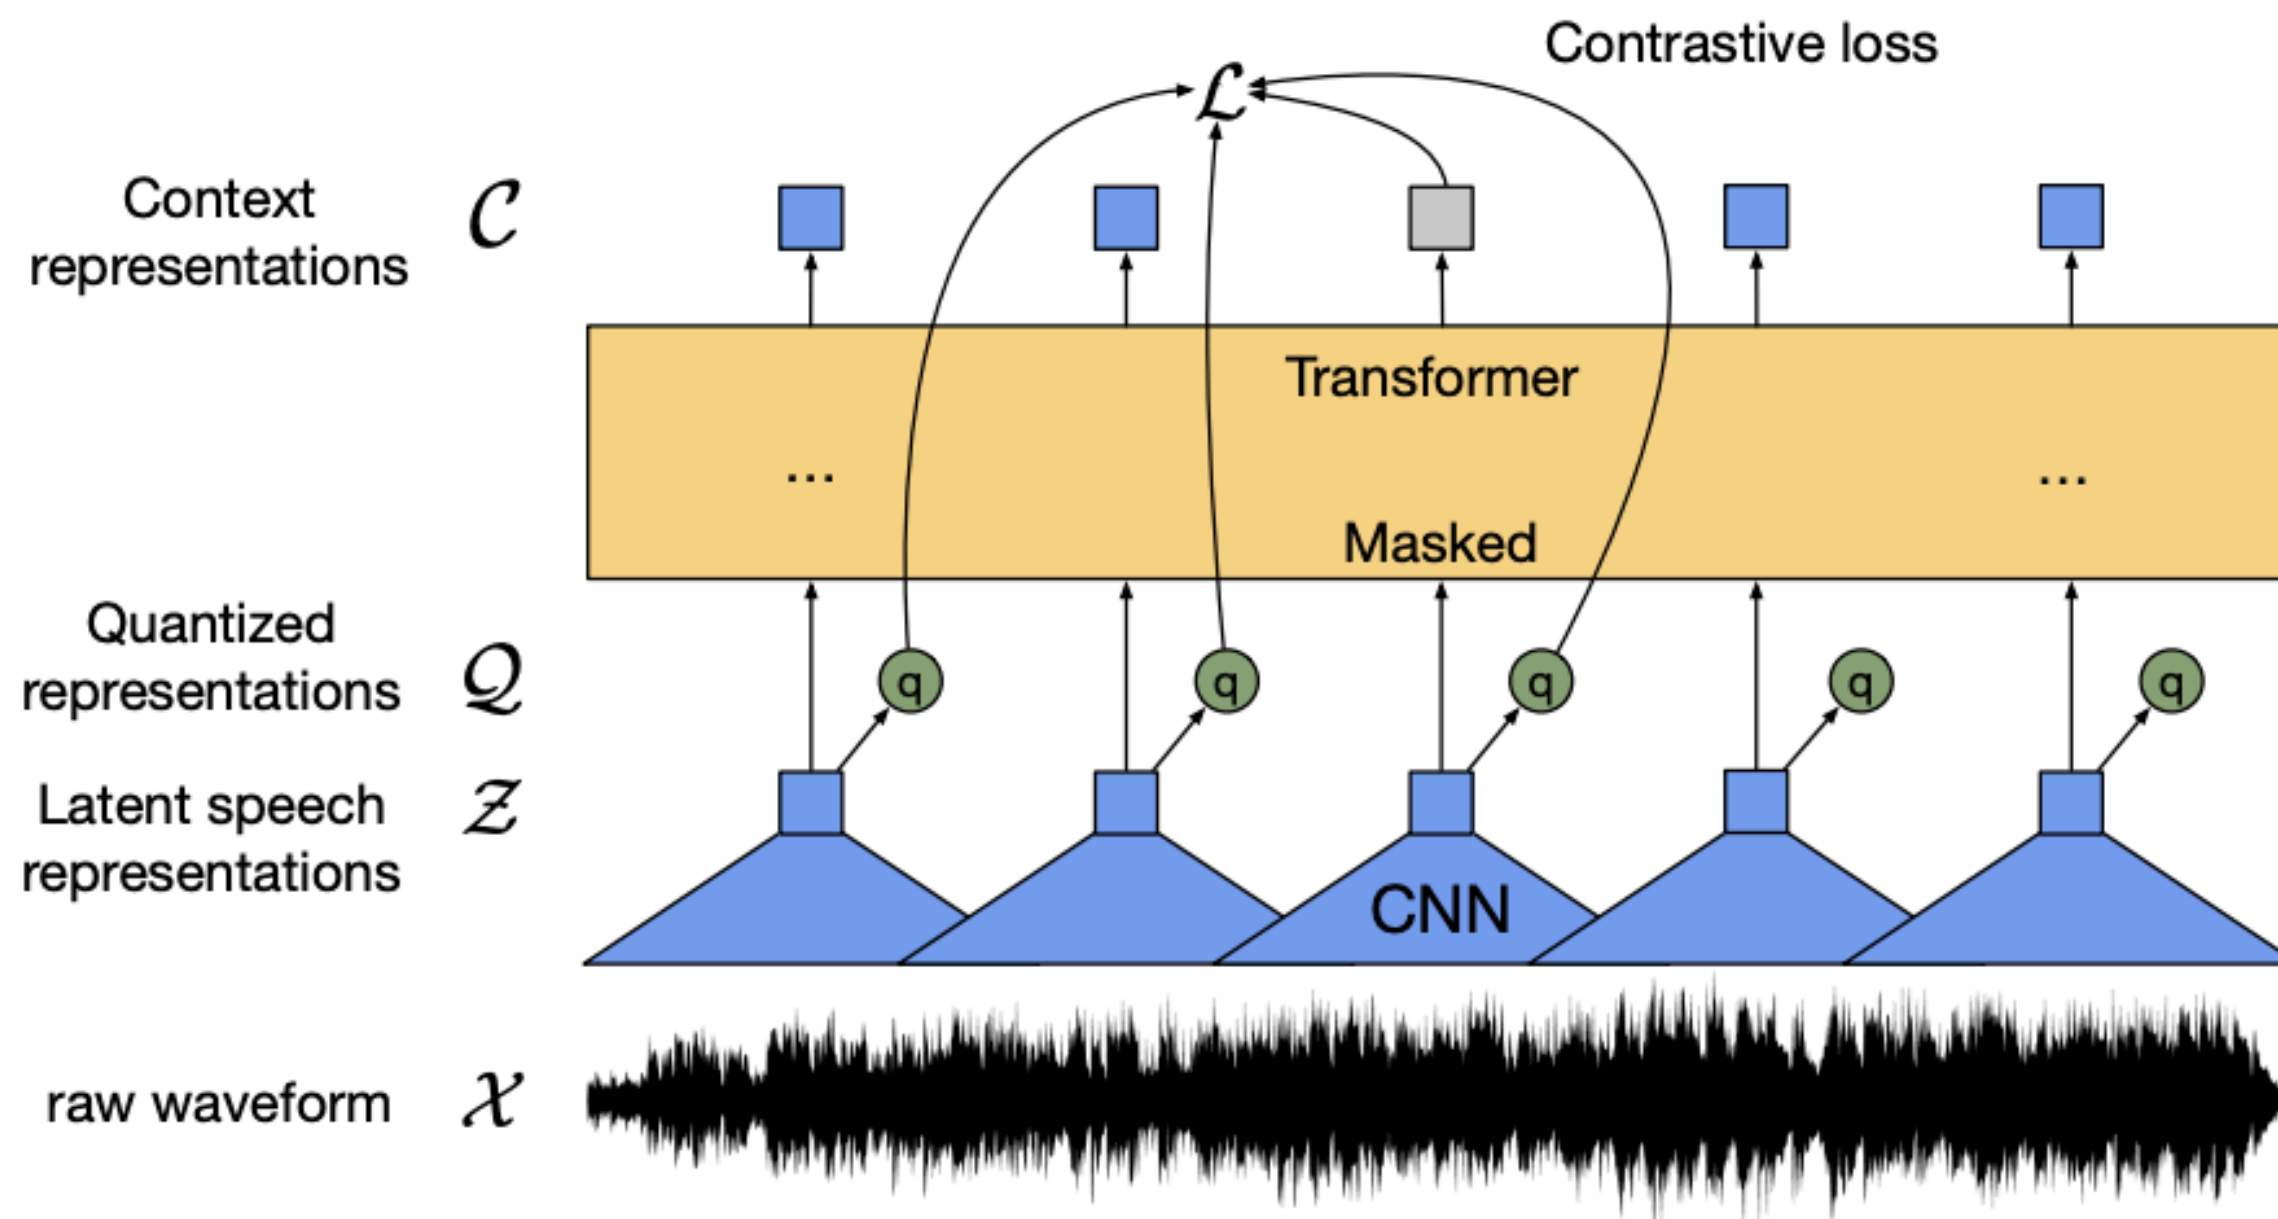


RAW AUDIO WAVEFORM



WAV2VEC 2.0

WAV2VEC 2.0



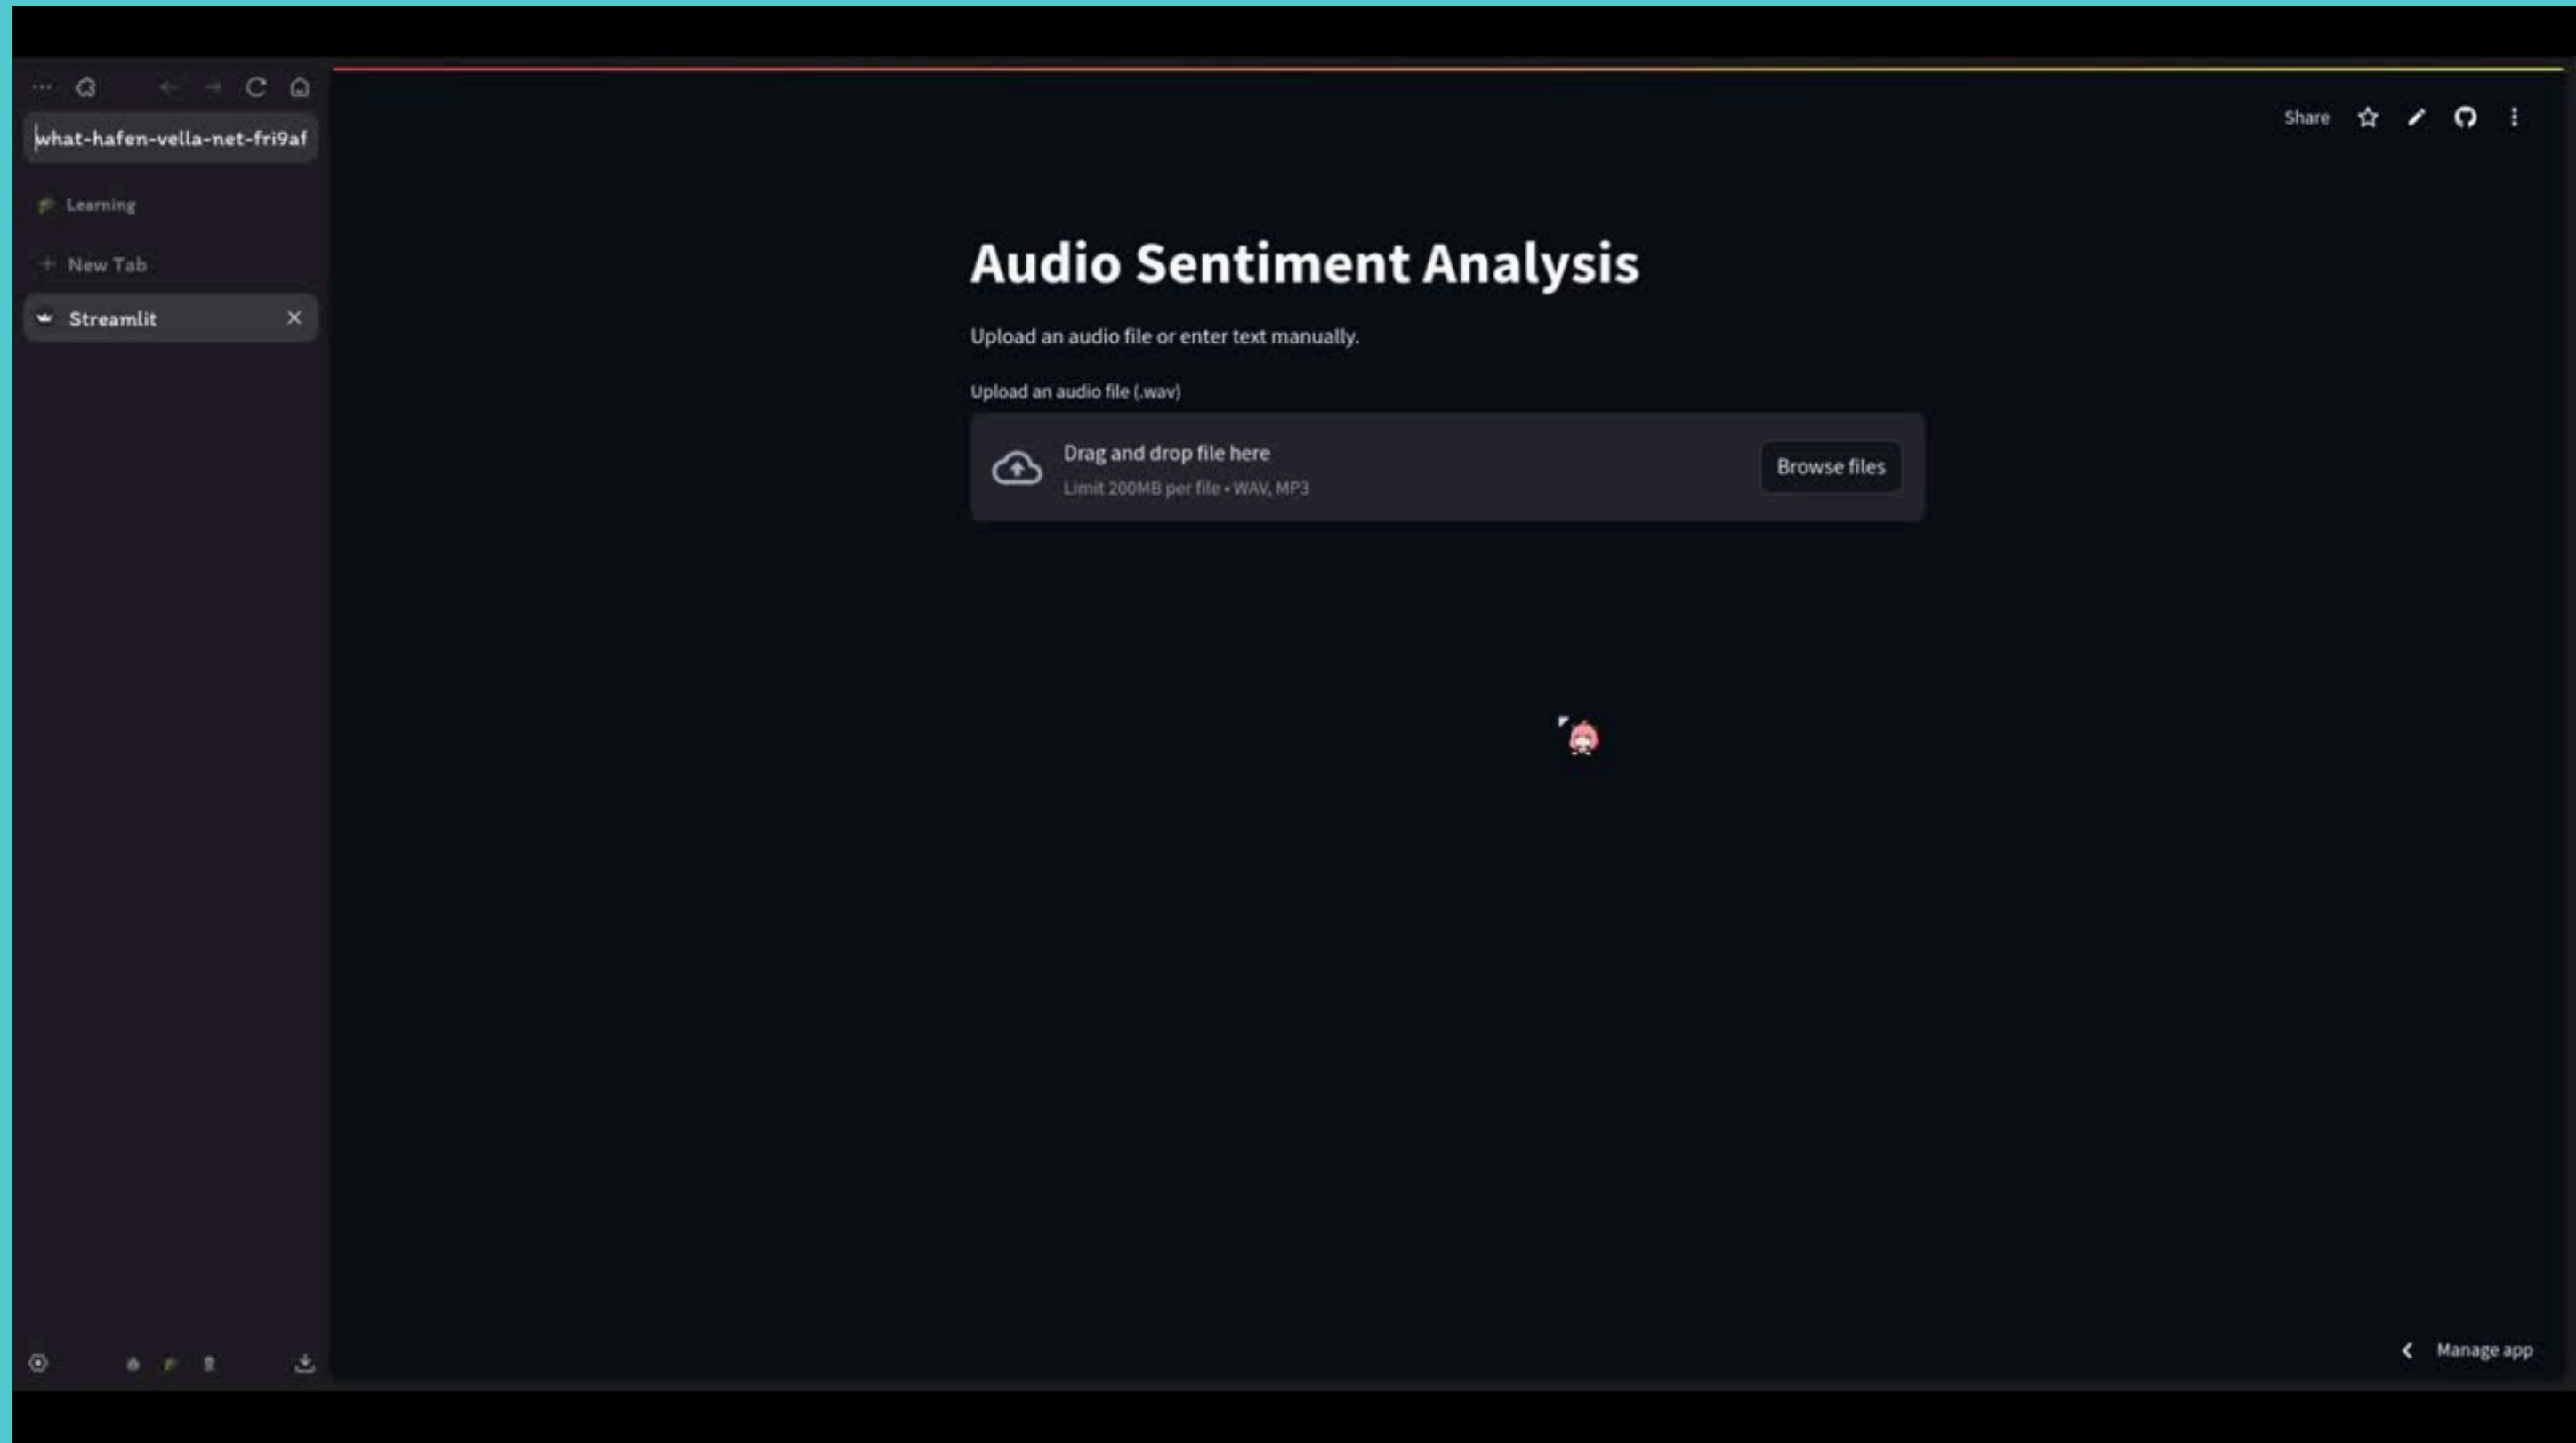
MODEL PERFORMANCE

| Metrics | CNN | ResNet50 | Wav2Vec 2.0 |
|---------------|--------|----------|-------------|
| Epoch | 25 | 10 | 5 |
| Accuracy | 0.9627 | 0.6562 | 0.9762 |
| F1 | 0.9625 | 0.6604 | 0.9762 |
| Cohen's Kappa | 0.9441 | 0.4843 | 0.9642 |

CONCLUSION

- For the text-based models, RoBERTa performed the best overall, reaching 92% accuracy and a strong F1 score of 0.91.
- On the audio side, Wav2Vec 2.0 delivered the highest performance among all models tested. It achieved 97.6% accuracy and a Cohen's Kappa of 0.964,
- both text and audio models have their strengths. Text-based models like RoBERTa are excellent when clean, accurate transcriptions are available. However, audio-based models, especially Wav2Vec 2.0, can recognize emotions directly from voice tone, making them more flexible in situations where the exact words are unclear or unimportant.

VIDEO PRESENTATION



SUSTAINABLE DEVELOPMENT GOAL

8 DECENT WORK AND ECONOMIC GROWTH



Promote sustained, inclusive and sustainable economic growth, full and productive employment and decent work for all.

8.5.1

AVERAGE HOURLY EARNINGS OF FEMALE AND MALE EMPLOYEES, BY OCCUPATION, AGE AND PERSONS WITH DISABILITIES

8.5.2

UNEMPLOYMENT RATE, BY SEX, AGE AND PERSONS WITH DISABILITIES



THANK YOU FOR LISTENING!

TEAM 13

**Dejoras, Dylan James
Villamor, Kurt Russel A.**